

Coursework 1 for M522 Numerical Analysis (Parallel Computing)

Pavel Buklemishev

February 21, 2025

- 1 Use OpenMP to speed up the code execution and demonstrate the difference from a serial execution
- 2 Considering one approach that uses a temporal array to store the results of multiplications before summation, can one use OpenMP reduction to reduce memory allocation?

Code:

```
1 #include "utils.h"
2 #include <omp.h>
3 #include <stdio.h>
4 #include <cmath>
5 using namespace std;
6 double compute_inner_nonparallel(long N, double * x, double *y){
7     double z = 0;
8
9     for(int i = 0; i<N; i++){
10         z += sqrt(x[i]*y[i]) ;
11     }
12
13     return z;
14 }
15 double compute_inner_reduction(long N, double* x, double* y) {
16
17     double z;
18
19     #pragma omp parallel for reduction(+:z)
20     for (long i = 0; i < N; i++){
21         z +=sqrt(x[i] * y[i]);}
22     return z;
23 }
24 double compute_inner_parallel(long N, double * x, double *y){
```

```

25 int num_threads = omp_get_max_threads();
26 double* sum_per_thread = (double*) malloc(num_threads* sizeof(
27   double));
28
29 #pragma omp parallel
30 {
31   int i_thread = omp_get_thread_num();
32   double i_sum = 0;
33   #pragma omp for
34   for(long i =0; i<N; i++){
35     i_sum += sqrt(x[i]*y[i]);
36   }
37
38   sum_per_thread[i_thread] = i_sum;
39 }
40 double z =0;
41 for(int i = 0; i< num_threads; i++)z += sum_per_thread[i];
42
43 free(sum_per_thread);
44 return z;
45 }
46
47 int main(int argc, char** argv) {
48
49   long N = 100000000;
50
51   double* x = (double*) malloc(N * sizeof(double));
52   double* y = (double*) malloc(N * sizeof(double));
53   for (long i = 0; i < N; i++) {
54     x[i] = i+1;
55     y[i] = 2.0 / (i+1);
56   }
57   double sum;
58   Timer t;
59   t.tic();
60
61   sum = compute_inner_nonparallel(N, x,y);
62
63   printf("res_nonparallel - malloc: %f time elapsed = %f\n",sum ,
64         t.toc());
65
66   t.tic();
67
68   sum = compute_inner_parallel(N, x,y);
69
70   printf("res_parallel - malloc: %f time elapsed = %f\n",sum , t .
71         toc());
72   t.tic();
73   sum = compute_inner_reduction(N, x, y);

```

```

73     printf("res_reduction: %f time elapsed = %f\n", sum, t.toc());
74
75     free(x);
76     free(y);
77     return 0;
78 }
```

Result:

```

1 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
2     pavel/parallel# mpiicpx -qopenmp inner_product.cpp
3 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
4     pavel/parallel# ./a.out
5 res nonparallel - malloc: 141421356.137423 time elapsed =
6     0.084770
7 res parallel - malloc: 141421356.238374 time elapsed =
8     0.051157
9 res reduction: 141421356.238374 time elapsed = 0.051621
```

Needed to add the sqrt for the demonstrational purposes.

3 Using optimization flags and/or different compilers, demonstrate the differences in the code execution times

g++

```

1 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
2     pavel/parallel# g++ -fopenmp inner_product.cpp
3 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
4     pavel/parallel# ./a.out
5 res nonparallel - malloc: 141421356.465761 time elapsed =
6     0.219311
7 res parallel - malloc: 141421356.238399 time elapsed = 0.048981
8 res reduction: 141421356.238399 time elapsed = 0.050593
```

g++ with o3 flag

```

1 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
2     pavel/parallel# g++ -fopenmp inner_product.cpp -O3
3 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
4     pavel/parallel# ./a.out
5 res nonparallel - malloc: 141421356.465761 time elapsed =
6     0.163962
7 res parallel - malloc: 141421356.238399 time elapsed = 0.045187
8 res reduction: 141421356.238399 time elapsed = 0.050634
```

icpx

```
1 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
  pavel/parallel# icpx -qopenmp inner_product.cpp
(base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
  pavel/parallel# ./a.out
res nonparallel - malloc: 141421356.137423 time elapsed =
  0.083285
res parallel - malloc: 141421356.238374 time elapsed = 0.053423
res reduction: 141421356.238374 time elapsed = 0.049398
```

icpx -o3

```
1 (base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
  pavel/parallel# icpx -qopenmp inner_product.cpp -O3
(base) root@pavel-Victus-by-HP-Gaming-Laptop-15-fa1xxx:/home/
  pavel/parallel# ./a.out
res nonparallel - malloc: 141421356.137423 time elapsed =
  0.083232
res parallel - malloc: 141421356.238374 time elapsed = 0.045548
res reduction: 141421356.238374 time elapsed = 0.043733
```

4 (Optional) Use a hybrid approach: MPI + OpenMP to further parallelise the code

Since my PC does not support MPI, I am attaching a draft of the code that implements the idea of MPI + OpenMP coupled parallelization as I can't debug it properly.

The idea:

1. Divide the vector into $N - 1$ equal parts and assign the remainder to the N th processor.
2. Apply the `compute_inner_reduction` function to each vector part to compute the partial inner sum.
3. Reduce the sums across processors by summing them to obtain the inner product of the entire vector.

```
1 #include "mpi.h"
2 #include <stdio.h>
3 using namespace std;
4 double compute_inner_reduction(long N, double* x, double* y) {
5
6     double z;
7
8     #pragma omp parallel for reduction(+:z)
9     for (long i = 0; i < N; i++) {
```

```

10         z +=(x[i] * y[i]);}
11     return z;
12 }
13
14 int main( int argc, char *argv[]
15 )
16 {
17     int rank, size;
18     MPI_Init( &argc, &argv );
19     MPI_Comm_rank( MPI_COMM_WORLD , &rank );
20     MPI_Comm_size( MPI_COMM_WORLD , &size );
21
22     long N = 100000000;
23     len_per_proc = N/size;
24     len_rest_proc= N%size;
25     double dotProduct;
26     if(rank==0){
27         double* x = (double*) malloc(N * sizeof(double));
28         double* y = (double*) malloc(N * sizeof(double));
29         for (long i = 0; i < N; i++) {
30             x[i] = i+1;
31             y[i] = 2.0 / (i+1);}
32     }
33     if(rank == size-1){
34         len_per_proc =len_rest_proc;
35     }
36
37     double sum=0;
38     double *x_i = (double*) malloc(len_per_proc * sizeof(double));
39     double *y_i = (double*) malloc(len_per_proc * sizeof(double));
40
41
42     MPI_Scatter(x, len_per_proc, MPI_DOUBLE, x_i, len_per_proc,
43                 MPI_DOUBLE, 0, MPI_COMM_WORLD);
44     MPI_Scatter(y, len_per_proc, MPI_DOUBLE, y_i, len_per_proc,
45                 MPI_DOUBLE, 0, MPI_COMM_WORLD);
46
47 //CODE THAT REALIZES VECTOR RESIDUE ALLOCATION TO THE LAST
48 // PROCESSOR .
49
50
51     sum = compute_inner_reduction(N, x_i, y_i);
52     MPI_Reduce(&sum, &dotProduct, 1, MPI_DOUBLE, MPI_SUM, 0,
53                MPI_COMM_WORLD);
54
55
56     cout<<dotProduct<<endl;
57
58     MPI_Finalize();
59     return 0;
60 }
```